

Learning from the expert: processing

CASE STUDY: SCHOOL BUDGETING WITH MACHINE LEARNING IN PYTHON




Peter Bull

Co-founder of DrivenData

Learning from the expert

- Text processing
- Statistical methods
- Computational efficiency



Quoc Le

100	100	1
ENTRIES	AVG #ENTRIES	VICTORIES

ABOUT QUOC

Northwestern University Masters in Predictive Analytics '14
Location: San Francisco, CA

1 COMPLETED COMPETITION

Box-Plots for Education
FINAL RANK: 1

Learning from the expert: text preprocessing

- NLP tricks for text data
 - Tokenize on punctuation to avoid hyphens, underscores, etc.
 - Include unigrams **and** bi-grams in the model to capture important information involving multiple tokens - e.g., "middle school"

N-grams and tokenization

```
vec = CountVectorizer(token_pattern=TOKENS_ALPHANUMERIC,  
                      ngram_range=(1, 2))
```

- Simple changes to `CountVectorizer`
- alphanumeric tokenization
- `ngram_range=(1, 2)`

Range of n-grams in scikit-learn

```
pl.fit(X_train, y_train)
```

```
Pipeline(steps=[('union', FeatureUnion(n_jobs=1,  
    transformer_list=[('numeric_features',  
Pipeline(steps=[('selector',  
FunctionTransformer(accept_sparse=False,  
    func=<function <lambda> at 0x11441f7b8>, pass_y=False,  
    validate=False)), ('imputer', Imputer(axis=0, copy=True,  
missing_valu...=None,  
    solver='liblinear',  
    tol=0.0001,  
    verbose=0, warm_start=False),  
    n_jobs=1)))]])
```

Range of n-grams in scikit-learn

```
holdout = pd.read_csv('HoldoutData.csv', index_col=0)
predictions = pl.predict_proba(holdout)
prediction_df = pd.DataFrame(columns=pd.get_dummies(
    df[LABELS]).columns, index=holdout.index,
    data=predictions)
prediction_df.to_csv('predictions.csv')
score = score_submission(pred_path='predictions.csv')
```

Let's practice!

CASE STUDY: SCHOOL BUDGETING WITH MACHINE LEARNING IN PYTHON

Learning from the expert: a stats trick

CASE STUDY: SCHOOL BUDGETING WITH MACHINE LEARNING IN PYTHON



Peter Bull

Co-founder of DrivenData

Learning from the expert: interaction terms

- Statistical tool that the winner used: interaction terms
- Example
 - English teacher for 2nd grade
 - 2nd grade - budget for English teacher
- Interaction terms mathematically describe when tokens appear together

Interaction terms: the math

$$\beta_1 x_1 + \beta_2 x_2 + \beta_3 (x_1 \times x_2)$$

Interaction terms: the math

$$\beta_1 x_1 + \beta_2 x_2 + \beta_3 (x_1 \times x_2)$$

Interaction terms: the math

$$\beta_1 x_1 + \beta_2 x_2 + \beta_3 (x_1 \times x_2)$$

X1	X2
0	1
1	1

Interaction terms: the math

$$\beta_1 x_1 + \beta_2 x_2 + \beta_3 (x_1 \times x_2)$$

X1	X2	X3
0	1	$X1 * X2 = 0 * 1 = 0$
1	1	$X1 * X2 = 1 * 1 = 1$

Adding interaction features with scikit-learn

```
from sklearn.preprocessing import PolynomialFeatures  
x
```


	x1	x2
a	0	1
b	1	1

```
interaction = PolynomialFeatures(degree=2,  
                                interaction_only=True,  
                                include_bias=False)  
  
interaction.fit_transform(x)
```

```
array([[ 0.,  1.,  0.],  
       [ 1.,  1.,  1.]])
```

A note about bias terms

- Bias term allows model to have non-zero y value when x value is zero



Sparse interaction features

```
SparseInteractions(degree=2).fit_transform(x).toarray()
```

```
array([[ 0.,  1.,  0.],  
       [ 1.,  1.,  1.]])
```

- The number of interaction terms grows exponentially
- Our vectorizer saves memory by using a sparse matrix
- `PolynomialFeatures` does not support sparse matrices
- We have provided `SparseInteractions` to work for this problem

Let's practice!

CASE STUDY: SCHOOL BUDGETING WITH MACHINE LEARNING IN PYTHON

Learning from the expert: the winning model

CASE STUDY: SCHOOL BUDGETING WITH MACHINE LEARNING IN PYTHON




Peter Bull

Co-founder of DrivenData

Learning from the expert: hashing trick

- Adding new features may cause enormous increase in array size
- Hashing is a way of increasing memory efficiency



- Hash function limits possible outputs, fixing array size

When to use the hashing trick

- Want to make array of features as small as possible
 - Dimensionality reduction
 - Particularly useful on large datasets
 - e.g., lots of text data!


Implementing the hashing trick in scikit-learn

```
from sklearn.feature_extraction.text import HashingVectorizer

vec = HashingVectorizer(norm=None,
                        non_negative=True,
                        token_pattern=TOKENS_ALPHANUMERIC,
                        ngram_range=(1, 2))
```

The model that won it all

- You now know all the expert moves to make on this dataset
 - NLP: Range of n-grams, punctuation tokenization
 - Stats: Interaction terms
 - Computation: Hashing trick
- What class of model was used?



Quoc Le

100	100	1
ENTRIES	AVG #ENTRIES	VICTORIES

ABOUT QUOC

Northwestern University Masters in Predictive Analytics '14
Location: San Francisco, CA

1 COMPLETED COMPETITION

Box-Plots for Education
FINAL RANK: 1 ✓

The model that won it all

- And the winning model was...
- Logistic regression!
 - Carefully create features
 - Easily implemented tricks
 - Favor simplicity over complexity and see how far it takes you!

Let's practice!

CASE STUDY: SCHOOL BUDGETING WITH MACHINE LEARNING IN PYTHON

Next steps and the social impact of your work

CASE STUDY: SCHOOL BUDGETING WITH MACHINE LEARNING IN PYTHON



Peter Bull

Co-founder of DrivenData

Can you do better?

- You've seen the flexibility of the pipeline steps
- Quickly test ways of improving your submission
 - NLP: Stemming, stop-word removal
 - Model: RandomForest, k-NN, Naïve Bayes
 - Numeric Preprocessing: Imputation strategies
 - Optimization: Grid search over pipeline objects
 - Experiment with new scikit-learn techniques
- Work with the full dataset at DrivenData!

Hundreds of hours saved

- Make schools more efficient by improving their budgeting decisions
- Saves hundreds of hours each year that humans spent labeling line items
- Can spend more time on the decisions that really matter

DrivenData: Data Science to save the world

- Other ways to use data science to have a social impact at www.drivendata.org
 - Improve your data science skills while helping meaningful organizations thrive
 - Win some cash prizes while you're at it!

Let's practice!

CASE STUDY: SCHOOL BUDGETING WITH MACHINE LEARNING IN PYTHON